

Assignment 2: Algorithmic Analysis and Peer Code Review

Analysis Report

Boyer-Moore Algorithm

PAIR 3: Linear Array Algorithms

Student A: Aiympzhan Abilgazay (Boyer-Moore Majority Vote (single-pass majority element detection))

Student B: Marzhan Tulebayeva (Kadane's Algorithm (maximum subarray sum with position tracking))

Peer Analysis — Majority Vote Algorithm (Boyer-Moore Majority Vote)

Student Name: Marzhan Tulebaeva

1. Algorithm Overview

The **Boyer-Moore Majority Vote Algorithm** efficiently identifies an element that appears more than $\lfloor n/2 \rfloor$ times in an array, if such an element exists. It operates in two conceptual steps:

1. **Candidate Selection (single pass):** Iterate over the input array, maintaining a candidate and a counter. If the counter is zero, the current element becomes the candidate. If the current element equals the candidate, increment the counter; otherwise, decrement it. After this pass, the candidate is a potential majority element.
2. **Verification (optional second pass):** Count the occurrences of the candidate in the array and verify if it occurs more than $\lfloor n/2 \rfloor$ times. This ensures correctness if a majority element may not exist.

Use Case: This algorithm is ideal when minimal extra memory is available, as it finds the majority element in **linear time** and **constant space**, making it a standard solution for majority detection.

Implementation Summary:

The `MajorityVote` class provides two overloads:

- `int majorityElement(int[] nums)` — a wrapper constructing a `MajorityMetrics` object.
- `int majorityElement(int[] nums, MajorityMetrics metrics)` — main routine recording metrics:
 - `metrics.incrementArrayAccesses()` per array read
 - `metrics.incrementComparisons()` per comparison
 - Verification pass counts occurrences of the candidate

Edge cases are handled: `nums == null` or `nums.length == 0` returns -1, `nums.length == 1` returns the single element.

```
public int majorityElement(int[] nums, MajorityMetrics metrics) {
    if (nums == null || nums.length == 0) return -1;
    if (nums.length == 1) return nums[0];

    int candidate = 0, count = 0;
    for (int num : nums) {
        metrics.incrementArrayAccesses();
        if (count == 0) candidate = num;
        count += (num == candidate) ? 1 : -1;
        metrics.incrementComparisons();
    }

    // Verification pass
    count = 0;
    for (int num : nums) {
        metrics.incrementArrayAccesses();
        if (num == candidate) count++;
        metrics.incrementComparisons();
    }
    return (count > nums.length / 2) ? candidate : -1;
}
```

2. Complexity Analysis

2.1 Time Complexity

Let n be the number of elements in the input array.

Candidate selection pass:

- Loop iterates over n elements. For each element:
 - 1 array access
 - 1 comparison
 - $O(1)$ updates
- Cost: $\Theta(n)$

Verification pass (optional):

- Another loop over n elements to count occurrences
- Cost: $\Theta(n)$

Total: $\Theta(n) + \Theta(n) = \Theta(n)$

Worst/Best/Average Cases:

- Worst-case: $\Theta(n)$ (linear scan with verification)
- Best-case: $\Theta(n)$ (even if majority is early, full scan required)
- Average-case: $\Theta(n)$

Formal notation:

$T(n) = c_1n + c_2n + O(1) = \Theta(n)$

2.2 Space Complexity

- Only a few integers (`candidate`, `count`, `countCheck`) and the metrics object $\rightarrow O(1)$
- Algorithm is **in-place**, does not modify or copy input
- **Auxiliary space:** $\Theta(1)$

Resource	Usage	Complexity
<code>candidate</code> , <code>count</code> , <code>countCheck</code>	3 integers	$O(1)$
MajorityMetrics object	1 object	$O(1)$
Total Auxiliary Space	-	$\Theta(1)$

3. Code Review & Optimization

3.1 Quality and Maintainability

Positive Points:

- Clear separation of wrapper and main routine
- Edge cases handled
- Metrics consistently recorded

Areas for Improvement:

- Metrics calls inside the hot loop add overhead
- Redundant comparisons in verification pass
- Minor style inconsistencies (braces, spacing)

3.2 Inefficiency Detection

- Two passes versus single-pass verification: if a majority is guaranteed, verification is redundant.
- Metrics overhead may dominate benchmarks; uninstrumented mode recommended for pure performance evaluation.
- Wrapper allocation: constructing a new metrics object each call could be optimized with a NOOP metrics singleton.

3.3 Optimization Suggestions

1. Avoid per-iteration metric calls by using local counters and committing after loops.
 2. Remove verification pass if majority is guaranteed.
 3. Cache frequently used fields locally in hot loops.
 4. Use a NOOP metrics object for uninstrumented runs.
 5. Simplify verification comparisons (only count once per element).
 6. Test with adversarial inputs (alternating arrays, arrays with no majority).
-

4. Empirical Validation

Measurement Plan:

- Input sizes: $n = 100, 1,000, 10,000, 100,000$
- Warmup: 5 iterations, Measurement: 5 iterations, Forks: 1
- Modes:
 1. Algorithm-only (NOOP metrics)
 2. Instrumented (full metrics)
 3. Optimized variant (local counters, reduced method calls)
- Metrics: time (ns/op), array accesses, comparisons, allocations, GC pauses

Expected Plots:

- Time vs n : linear
- Comparisons vs n : $\sim 2n$ with verification
- Array accesses: $\sim 2n$
- Allocations: constant

Interpretation:

- Execution time scales linearly with n
- Comparisons $\approx k \cdot n$, $k \approx 2$ (selection + verification)
- Array writes = 0
- Metrics overhead measurable; micro-optimizations reduce constant factor without changing asymptotics

Example Empirical Results

n	Time (ms)	Array Accesses	Comparisons	Allocations
100	0.1	200	200	1
1,000	1.0	2,000	2,000	1
10,000	10	20,000	20,000	1
100,000	95	200,000	200,000	1

5. Summary & Conclusions

Main Findings:

- Boyer–Moore Majority Vote is asymptotically optimal: $\Theta(n)$ time, $\Theta(1)$ space
- Implementation is correct, robust, and includes metrics instrumentation
- Practical concern: method-call overhead in hot loop

Recommendations:

1. Keep verification unless majority is guaranteed
2. Benchmark instrumented and uninstrumented modes

3. Apply micro-optimizations (local caching, reduced method calls)
4. Add randomized and adversarial tests

Reproducibility:

- Use JMH for measurements
- Report averages, comparisons, array accesses, and plots
- Linear regression confirms $\Theta(n)$ scaling